



Nonprofit Financial Records

A Guide to the Dataset

August 2026 · v1.2.1

Part 1: Substantive Insights

1. Overview

This dataset provides longitudinal financial data — balance sheet items, profit and loss aggregates, employee counts, and CIC34 community-impact narratives — for nonprofit companies registered at Companies House. It was produced by the [UK Third and Civil Society Sector Database](https://uk-third-sector-database.github.io/) (<https://uk-third-sector-database.github.io/>) project, which collects, processes, and links public administrative data on civil society organisations across the United Kingdom.

The data is drawn from the annual accounts that nonprofit companies file at Companies House. Two complementary sources are used: structured XBRL filings (machine-readable accounts) and PDF accounts processed through a structured-output LLM extraction pipeline. The merged dataset gives broader coverage than either source alone, particularly for smaller filings such as Community Interest Company (CIC) abridged accounts.

The dataset spans financial years from 1981 to 2026 and contains 1,342,198 financial-year records linked to 217,352 nonprofit organisations. Each row represents one company's accounts for a single financial year.

1,342,198 Financial Records	217,352 Linked Nonprofit Organisations
2 Account Sources	1981–2026 Financial Years

2. What are Nonprofit Financial Records?

Companies registered at Companies House are required to file annual accounts. The level of detail filed varies by company size — micro-entities and small companies file abridged or filleted accounts, while medium and large companies file full statutory accounts. CICs additionally file a CIC34 community-impact report alongside their accounts.

XBRL Accounts (Companies House monthly extracts)

Companies House publishes monthly bulk extracts of accounts filed in machine-readable XBRL format. The extracts contain structured tagged values for balance sheet items, profit and loss aggregates, and employee counts. These are downloaded from the [Companies House accounts download site](https://download.companieshouse.gov.uk/en_monthlyaccountsdata.html) (https://download.companieshouse.gov.uk/en_monthlyaccountsdata.html).

PDF Accounts (LLM-extracted)

For older filings and CIC abridged accounts that pre-date or were filed outside the XBRL stream, the project extracts financial line items from the original PDF accounts using a structured-output LLM pipeline. The extracted values are mapped onto the same column schema used for XBRL, so the merged dataset has a uniform shape regardless of source.

CIC34 Community Interest Reports

Community Interest Companies file a CIC34 alongside their accounts. The CIC34 contains four narrative sections: the company’s activities and impact, stakeholder consultation, directors’ remuneration, and any asset transfers. These narratives are extracted from the same PDF filings.

3. Dataset Contents

The published file has 66 columns, organised into the groups below. Each row represents one company-year record. Where a value is missing, the cell is left blank.

Identifiers and Time Period

Field	Description	Type
uid	The organisation’s identifier in the project Spine, and the key that links every row for the same organisation across years and sources (for example <code>GB-COH-08411754</code> , <code>GB-CHC-1000019</code>)	Text
coyno	Companies House registered number (8-digit, zero-padded)	Text
entity_current_legal_name	Company name as filed	Text
fy	Financial year (calendar year of FYE)	Numeric
fys / fye	Financial year start / end dates	Date
balance_sheet_date	The date the balance sheet in the filing was drawn up to	Date

Aggregate Financials

Field	Description
turnover_gross_operating_revenue	Total turnover / operating income
operating_profit_loss	Operating profit or loss
profit_loss_for_period	Profit or loss for the period

Balance Sheet

Balance-sheet line items including `tangible_fixed_assets` , `debtors` , `cash_bank_in_hand` , `current_assets` , `creditors_due_within_one_year` , `creditors_due_after_one_year` , `net_current_assets_liabilities` , `total_assets_less_current_liabilities` , `net_assets_liabilities_including_pension_asset_liability` , `called_up_share_capital` , `profit_loss_account_reserve` , and `shareholder_funds` .

Profit & Loss Detail

Detailed P&L items including `cost_sales` , `gross_profit_loss` , `administrative_expenses` , `raw_materials_consumables` , `staff_costs` , `depreciation_other_amounts_written_off_tangible_intangible_fixed_assets` , `other_operating_charges_format2` , `profit_loss_on_ordinary_activities_before_tax` , `tax_on_profit_or_loss_on_ordinary_activities` , `wages_salaries` , `dividends_paid` , and `government_grant_income` .

Staff

`average_number_employees_during_period` : average number of persons employed during the financial year.

CIC34 Narratives (CICs only)

Four free-text fields extracted from the CIC34 form: `cic34_activities_impact` , `cic34_stakeholder_consultation` , `cic34_directors_remuneration` , `cic34_asset_transfer` .

Where Each Row and Figure Came From

These columns describe the origin of the row and of the individual figures on it. They are a record of how the data was assembled, not measurements of the organisation. Use them to judge how much weight to put on a value.

Field	Description
source_dataset	Which stream the row was built from: <code>xbrl</code> for a machine-readable Companies House filing, or <code>pdf_extraction</code> for a PDF filing read by the extraction pipeline.
source_file , file_type , taxonomy	The archive the row was read from, the filing format, and the accounting taxonomy the original document was tagged against.
from_prior_year	TRUE when the whole row was recovered from the comparative (prior-year) figures printed in the <i>following</i> year's filing, because no filing of the organisation's own could be found for that year. FALSE for an ordinary row taken from a filing made for that year.
prior_source_fy	For a recovered row, the financial year of that later filing.
prior_source_file	For a recovered row, the archive that later filing was read from.
prior_backfilled_fields	On an ordinary row, a comma-separated list of the fields whose cells were blank in the organisation's own filing and were filled from the comparative figures in the following year's filing. Blank when nothing was filled in this way.

Field	Description
prior_backfill_source_fy	The financial year of the later filing that supplied those filled cells.
prior_backfill_source_file	The archive that filing was read from.
financial_cell_sources	A cell-level record of which source supplied the financial figures on the row. On a row built from an XBRL filing it names, comma-separated, the financial fields whose value was taken from the PDF extraction because the XBRL cell was blank, and is itself blank when every figure came from the XBRL filing. On a row with no XBRL filing behind it, it reads <code>pdf_extraction</code> , meaning every figure was extracted from the PDF.
derivation_source	The origin of any value worked out arithmetically from other cells rather than read from a filing. No such derivations are enabled in this release (see Section 6), so this column is blank throughout.

Organisation Attributes from the Spine

These columns are attached to each row from the project Spine, the register of organisations described in Annex A2.

Field	Description
uid	The organisation's Spine identifier, as above. It is the join key between this dataset and any other product built on the Spine.
normalised_name	The organisation's name after standardising case, punctuation, and common abbreviations. It exists to make names comparable between registers; the name exactly as filed is in <code>entity_current_legal_name</code> .
is_cic	TRUE when the Spine records the organisation as a Community Interest Company.
csotype	The organisation's broad civil-society type as recorded in the Spine: <code>CIC</code> , <code>Charity</code> , <code>Co-operative / Mutual</code> , or <code>Other</code> .

4. Coverage & Completeness

The combined dataset draws from XBRL and PDF sources. Coverage varies by financial year, by company size, and by individual line item. The charts and tables below summarise where the data is fullest and where it is sparsest.

Coverage is reported as a diagnostic description of this release. It is not a guarantee that every filing or line item is present, and the release policy does not impose arbitrary absolute coverage floors.

The coverage rate below compares the organisations with accounts data in a year against the organisations that the release Spine — the same Spine this release was built from (Annex A2) — records as active at the end of that year, that is registered on or before 31 December and not yet removed.

Records by Source and Year




Figure: Records by source and year. Earlier years are XBRL-dominated; recent CIC filings come predominantly via the PDF extraction stream.

Source Overlap




Figure: Most records are uniquely sourced; the XBRL stream is preferred where both are available.

Variable Availability




Figure: Fill rates for the financial columns across all records.

Headline Coverage

Dataset Summary	
Combined accounts dataset headline statistics	
Metric	Value
Total Records	1,342,198
Unique Organisations	217,352
Unique CICs	43,405
Columns in the Published File	66
Earliest Financial Year	1981
Latest Financial Year	2026
XBRL Records	1,194,620
PDF Extraction Records	147,578

Year-by-Year

Coverage by Financial Year			
Unique organisations with accounts data per year			
Financial Year	Unique Orgs	Total Active Orgs	Coverage Rate
2010	14,884	296,033	5.0%
2011	27,254	302,320	9.0%
2012	41,325	291,847	14.2%
2013	52,853	298,929	17.7%
2014	62,419	304,716	20.5%

Coverage by Financial Year

Unique organisations with accounts data per year

Financial Year	Unique Orgs	Total Active Orgs	Coverage Rate
2015	73,504	312,785	23.5%
2016	81,096	320,250	25.3%
2017	85,278	323,495	26.4%
2018	91,498	324,720	28.2%
2019	98,304	330,993	29.7%
2020	105,943	343,987	30.8%
2021	114,042	350,916	32.5%
2022	116,785	357,835	32.6%
2023	123,068	365,147	33.7%
2024	130,534	372,905	35.0%
2025	99,732	381,101	26.2%

5. What Can You Learn?

The dataset supports questions about the size, financial health, and activity of the UK nonprofit company sector. A non-exhaustive list of uses:

Sector Growth Trends




Figure: CICs with reported accounts and median net assets by financial year.

Financial Profile




Figure: Distribution of five high-coverage balance-sheet fields across the sector.

Employment Dynamics




Figure: Total employees and reporting organisations by financial year for CIC accounts.

CIC34 Narratives

The CIC34 narrative fields support qualitative analysis of community impact and stakeholder engagement. Coverage rates by field:

CIC34 Narrative Field Coverage			
Coverage rates for CIC-specific narrative fields			
Field	Non-Missing	Total CIC Records	Coverage Rate
Activities & Impact	176,318	204,437	86.2%
Stakeholder Consultation	174,002	204,437	85.1%
Directors Remuneration	172,304	204,437	84.3%
Asset Transfer	168,803	204,437	82.6%

Director Remuneration

Director remuneration patterns can be inferred from the CIC34 directors’ remuneration narrative. Among CICs reporting remuneration: coverage rate 84%, median statement length 33 characters.

6. Limitations & Caveats

Accepted release limitations. The counts in this section come from the passing, hash-bound release manifest. They describe known omissions; they have not been converted into apparently observed values.

Reporting Lag

Accounts are typically filed 6–12 months after the financial year ends. The most recent financial years will therefore be incomplete.

Filing Thresholds

Many CICs file as micro-entities and are exempt from disclosing detailed P&L information. Coverage of fields such as `turnover_gross_operating_revenue` and `staff_costs` is materially lower than balance-sheet coverage as a result.

Variable Financial Year Periods

Financial years vary in length, particularly for newly incorporated or dormant companies. Compare values per-year cautiously when the FY length differs from twelve months.

Empty PDF Extractions

The conversion quarantined 10,717 source filings whose extraction contained no usable values: 10,243 were not classified as dormant and 474 were classified as dormant. 1 local source PDF was unavailable for review. These cases remain in the release evidence rather than being filled with inferred values.

All-Zero PDF Rows

Among current PDF rows that are not explicitly dormant (including rows where dormancy is unknown), 8,233 of 124,675 have all twelve balance-sheet fields recorded as zero or blank (6.604%). This exceeds the unchanged 2.000% default diagnostic threshold. It is an accepted release exception and an extraction pattern, not evidence that every underlying company had a zero balance sheet.

PDF stratum	All-zero rows	Rows	Share
Current, explicitly non-dormant	5,069	84,856	5.974%
Current, dormancy unknown	3,164	39,819	7.946%
Current, explicitly dormant	8,909	14,512	61.391%
Prior comparative	2,506	8,391	29.865%

Conflicting PDF Candidates

The converter recorded 81,734 material conflict events, representing 46,048 distinct source-period-field cells omitted from the converted PDF data. It emitted 0 disputed values. When direct candidates disagree, no winner is selected: the published cell is blank unless a non-conflicting preferred source supplies it. In the assembled final population, 28,724 affected cells remain blank and 10,792 are safely populated by a separate, non-conflicting auditable source.

Disabled Arithmetic Derivations

The pipeline does not manufacture additional coverage using the following identities because their PDF-source confirmation checks failed:

- $\text{creditors_due_within_one_year} = \text{current_assets} - \text{net_current_assets_liabilities}$ — PDF identity check failed.
- $\text{current_assets} = \text{net_current_assets_liabilities} + \text{creditors_due_within_one_year}$ — PDF identity check failed.

The release manifest does not record this list, so the two identities above are the guidance generator's documented default. Unlike the counts elsewhere in this section, they have not been read from the hash-bound manifest.

How to Interpret Blanks

A blank can mean that the filing did not disclose the item, extraction produced no usable value, or conflicting candidates were deliberately suppressed. Treat blanks as missing information, not as zero. The `source_dataset` column can be used to distinguish XBRL and PDF provenance.

Spine Filtering

Only nonprofit companies present in the project Spine (the deduplicated register of UK third-sector and civil-society organisations) are included. Profit-making companies and other Companies House registrants are excluded.

What is NOT in the Data

The dataset does not include cash-flow statements, notes to the accounts in full, contingent liabilities, related-party transactions, or auditor qualifications. For these, consult the original PDFs at Companies House.

7. Citation & Licence

The dataset is released under the [Creative Commons Attribution 4.0 International \(CC BY 4.0\)](https://creativecommons.org/licenses/by/4.0/) (<https://creativecommons.org/licenses/by/4.0/>) licence. You are free to share and adapt the data for any purpose, including commercially, provided you give appropriate credit.

Suggested Citation

McDonnell, D. et al. (2026). Nonprofit Financial Records.
UK Third and Civil Society Sector Database.
Available at: <https://uk-third-sector-database.github.io/data/>
Licensed under CC BY 4.0.

8. Changelog

v1.2.1 — August 2026

- Point release: no financial value changed.
- Removed the `source_filing_date` column. It had been populated on about 706,000 XBRL rows with the following filing's accounting date even though nothing was recovered on those rows.
- Removed 24 undocumented `*__donor_end` working columns (91 columns down to 66).
- Coverage-rate denominators now use the release Spine v1.2 rather than the March 2026 spine.
- Documented the provenance columns in Section 3.

v1.2 — July 2026

- Rebuilt the XBRL and converted-PDF inputs and refreshed all figures and tables from the exact release candidate.
- Applied source-backed scalar/date/employee corrections and a pinned mapping no-regression audit.
- Removed arbitrary default coverage floors; coverage remains a diagnostic.
- Documented accepted empty-extraction, all-zero-row, and suppressed conflict limitations from the passing release manifest.
- Kept both unconfirmed arithmetic coverage derivations disabled.

v1.1 — May 2026

- Incorporated ~42,000 newly extracted CIC PDF accounts.
- Added `csotype` column (CIC / Charity / Co-operative / Mutual / Other).
- Refreshed coverage statistics throughout the guidance.

v1.0 — February 2026

- Initial public release of the merged XBRL + PDF combined dataset.

Part 2: Technical Annex

This annex provides the full technical detail on data sources, processing steps, coverage statistics, and known limitations for the combined accounts dataset.

A1. Pipeline Architecture

The dataset is produced by a multi-stage pipeline that combines two source streams into a single deduplicated CSV.

Stream 1: XBRL Accounts (monthly bulk extracts)

Companies House publishes monthly ZIP archives of XBRL-tagged accounts. Each archive is downloaded, parsed using `stream-read-xbrl`, and appended to a per-month CSV with a uniform schema. The combined extract represents the 'XBRL' stream.

Stream 2: PDF Accounts (LLM-extracted)

For each PDF account filing, the pipeline (a) renders pages to JPEG images, (b) submits the image set to OpenAI's Batch API with a JSON-schema response format that prompts a structured extraction of line items, balance-sheet date, and CIC34 narratives, and (c) collects the JSON outputs to disk under `data/output/accounts-extractions/json-schema/{coyno}/{txn}.json`.

Stream Convergence: XBRL Format Conversion

The PDF JSON extractions are converted into the same column schema as the XBRL stream using an approved label mapping plus a Jaccard-similarity fallback for unmapped labels. Sign-aware columns (creditors, expenses) are normalised to positive magnitudes.

Merge and Provenance

The XBRL and PDF-converted CSVs are merged on `(coyno, fy)`. XBRL is preferred where both sources cover the same record. A `source_dataset` column records each row's provenance.

A2. Source Data

Companies House — XBRL Accounts

The official monthly bulk download of accounts in XBRL format is available at download.companieshouse.gov.uk/en_monthlyaccountsdata.html (https://download.companieshouse.gov.uk/en_monthlyaccountsdata.html). Historic archives are at [historicmonthlyaccountsdata.html](https://download.companieshouse.gov.uk/historicmonthlyaccountsdata.html) (<https://download.companieshouse.gov.uk/historicmonthlyaccountsdata.html>).

Companies House — PDF Accounts

For filings outside the XBRL stream (older accounts, CIC abridged filings, etc.), original PDFs are obtained via the Companies House document API and stored locally for extraction.

Companies House — CIC34 Forms

Community Interest Reports are filed as part of the same PDF accounts package for CICs and are extracted alongside the financial line items.

Spine

The hash-bound release Spine (`code/release-v1.2-Jul2026/TSCS_spine.spine.csv`) is the deduplicated register of UK third-sector and civil-society organisations. It is the inclusion list for the dataset and supplies `uid` , `normalised_name` , `is_cic` , and the derived `csotype` .

A3. Deduplication & Validation

Within-Source Deduplication

For each (coyno, fy) pair, the pipeline prefers current-year filings over prior-year companion rows. Same-source duplicate rows (for example, balance-sheet and P&L sections of the same filing) may contribute non-conflicting values to one row. It does not simply take the first non-null value: distinct direct candidates for the same source-period-field cell are treated as a material conflict, logged, and omitted. This release logged 81,734 such events and omitted 46,048 disputed cells and emitted 0 unsafe disputed values. In the assembled final population, 28,724 of the affected cells remain blank and 10,792 are populated from a separate, non-conflicting auditable source.

Cross-Source Deduplication

Where both XBRL and PDF sources cover the same (coyno, fy), the XBRL row is preferred. The PDF row is dropped and not retained as a duplicate.

Prior-Period Recovery

Accounts print the previous year's figures alongside the current year's. Where an organisation has no filing of its own for a year, the pipeline can recover that year from the comparatives in the following year's filing; the row is flagged `from_prior_year` and the donor filing is named in `prior_source_fy` and `prior_source_file`. On an ordinary row, individual blank cells may be filled the same way; when that happens the fields filled are listed in `prior_backfilled_fields`, with the donor filing named in `prior_backfill_source_fy` and `prior_backfill_source_file`. A value that is already present is never overwritten.

The recovery step also uses internal working columns to carry the donor period's own end date while it matches filings. Those working columns are not part of the published dataset: they were removed in v1.2.1 (see Section 8), together with a filing-date column that had been stamped on every row with a later filing, whether or not anything was recovered from it. Nothing was published as a submission date, and no financial value depended on either column.

Date Validation

`fy` is normalised to integer-string. Where `fy` is blank but `balance_sheet_date` is populated, `fy` is derived (with a one-year offset for prior-year companion rows). Missing `fye` is imputed from the company's most common financial-year-end month-day.

Sign Normalisation

XBRL stores expenses and liabilities as positive magnitudes; PDF extractions sometimes preserve negative accounting signs. The pipeline applies `abs()` to nine sign-aware columns (`creditors_*`, `cost_sales`, `administrative_expenses`, etc.) so both sources use the same convention.

Arithmetic Derivations

Potential balance-sheet identities are disabled unless they pass source-stratified confirmation checks. The two candidate derivations in this release remain disabled because their PDF identity checks failed.

Employee-Count Guardrails

Legitimate non-negative fractional employee counts are preserved. Values tagged as currency or percentages, negative values, and counts above 100,000 are rejected. The same guard applies to direct, recovered, and prior-year candidates; if a current employee fact is rejected, the pipeline does not silently replace it with a prior-period value.

Spine Filtering

Only company numbers present in the Spine are retained. A left-join on `uid` populates `is_cic` and the derived `csotype`.

A4. Reproducibility

Code

The end-to-end pipeline lives in `code/companies-house/`:

- `pdf_accounts_extraction/openai-api/` — PDF-to-JSON extraction (rendering, batch submission, batch processing).
- `pdf_accounts_extraction/xbrl_format_conversion/` — JSON-to-XBRL column schema conversion, label mapping, dataset merge, and public-zip publishing.
- `reporting/guidance/` — this guidance document generator.

Dependencies

Managed via the [uv \(https://github.com/astral-sh/uv\)](https://github.com/astral-sh/uv) package manager. Run `python code/companies-house/pdf_accounts_extraction/xbrl_format_conversion/dependencies.py` to install conversion-pipeline dependencies; analogous scripts cover the other modules.

Running the Pipeline

The root ordinary build now names its two formerly hidden final stages. `convert-pdf` creates one explicit, versioned PDF CSV and `merge-final` consumes that exact path to build the actual final financial dataset. `combine-full` remains the XBRL-only intermediate; release gating and publication remain separate.

```
cd code/companies-house
python run_pipeline.py convert-pdf --xbrl-csv XBRL.csv --output CONVERTED_PDF.csv --diagnostics-dir PDF_DIAGNOSTICS
python run_pipeline.py merge-final --xbrl-csv XBRL.csv --pdf-csv CONVERTED_PDF.csv --spine-csv SPINE.csv --matches-csv MATCHES.csv --output COMBINED.csv

cd pdf_accounts_extraction/xbrl_format_conversion
python run_pipeline.py release 2026-08-16-spine-v1.2.1 --converted-csv CONVERTED_PDF.csv --conversion-report CONVERSION_REPORT.csv --check-b-audit CHECK_B_AUDIT.json --xbrl-csv XBRL.csv --spine-csv SPINE.csv --matches-csv MATCHES.csv --output-dir OUTPUT_DIR --accepted-limitations-policy POLICY.json

python ../../reporting/guidance/generate_guidance.py --combined-csv cso-spine-nonprofit-company-accounts-combined-2026-08-16-spine-v1.2.1.csv --pdf-csv cic-accounts-extractions-xbrl-format-2026-07-22-spine-v1.2-final-r3.csv --release-manifest manifest.json --release-date 2026-08-16 --release-version v1.2.1

python run_pipeline.py finalize-release --manifest manifest.json --guidance-pdf GUIDANCE.pdf --guidance-html GUIDANCE.html --licence-template LICENCE.txt --release-repo RELEASE_REPO --website-guidance WEBSITE_GUIDANCE --public-zip-name PUBLIC.zip --archive-date 2026-08-16
```

Guidance generation verifies the SHA-256 of both CSV inputs against the passing candidate manifest. Finalization does not rebuild that immutable candidate: it verifies the candidate and retained limitation evidence, then binds the explicit guidance pair, licence, and destinations by SHA-256 in a separate publication manifest before promotion. A lexical “latest file” lookup is not part of the release procedure.

Source Code

The release repository is at github.com/uk-third-sector-database/tso-database-builder (<https://github.com/uk-third-sector-database/tso-database-builder>).
